

10707

Deep Learning: Spring 2020

Andrej Risteski

Machine Learning Department

Lecture 2:

Benefits of depth

Recap

Recall from previous lecture:

(1): Neural networks are **universal approximators**: given any (reasonably nice) function $f: \mathbb{R}^d \rightarrow \mathbb{R}$, a **shallow** (3-layer) neural network with $\sim \left(\frac{1}{\epsilon}\right)^d$ neurons can approximate it to within ϵ error.

(2): Neural networks are **circumvent** curse of dimensionality for functions $f: \mathbb{R}^d \rightarrow \mathbb{R}$ with appropriately decaying Fourier coefficients: **shallow** (-layer) neural networks with $\sim \left(\frac{1}{\epsilon}\right)$ neurons can approximate them to within ϵ error.

Is there any point to depth?
Are deeper networks more powerful?

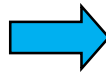
Part of the deep learning story



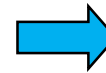
Object
detection



Image

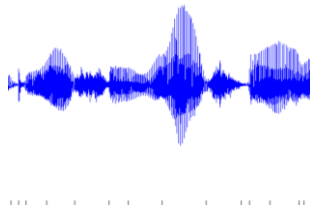


vision features

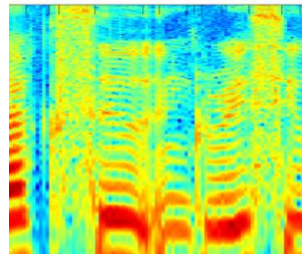
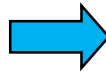


Recognition

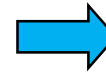
Audio
classification



Audio

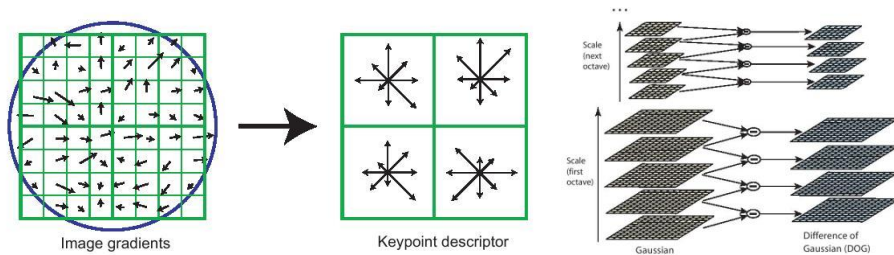


audio features

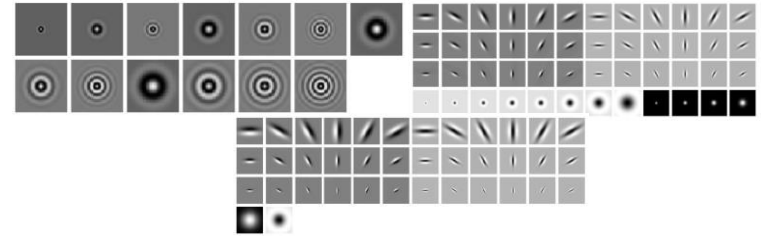


Speaker
identification

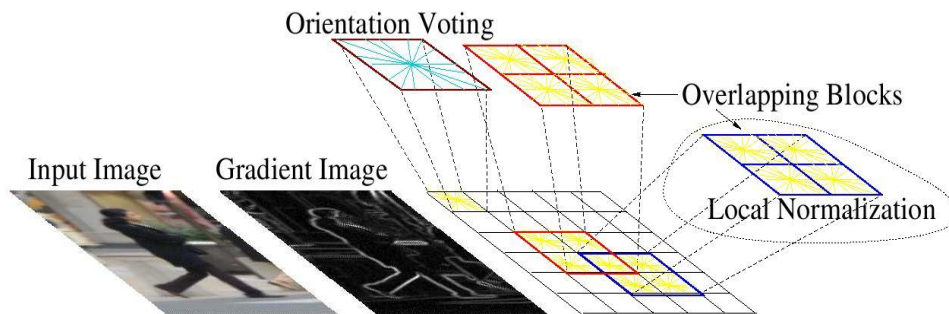
Old school: hand-craft features



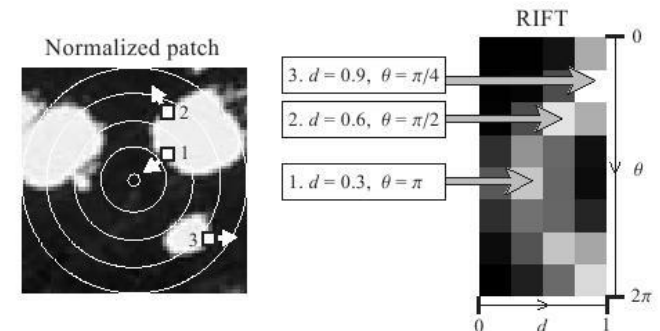
SIFT



Textons

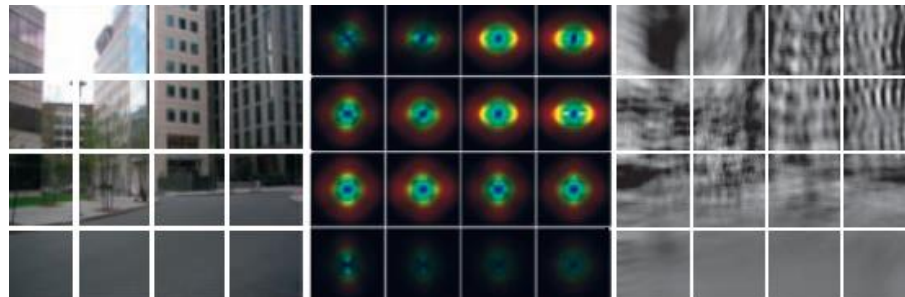


HoG

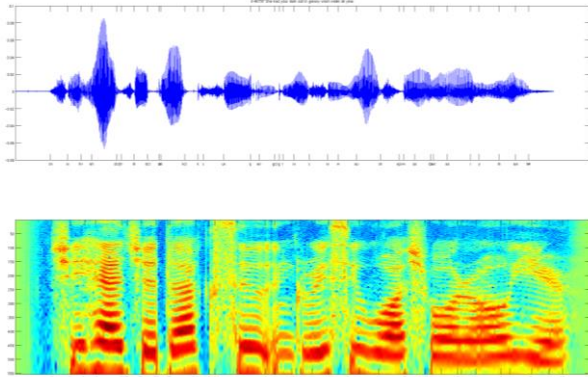


RIFT

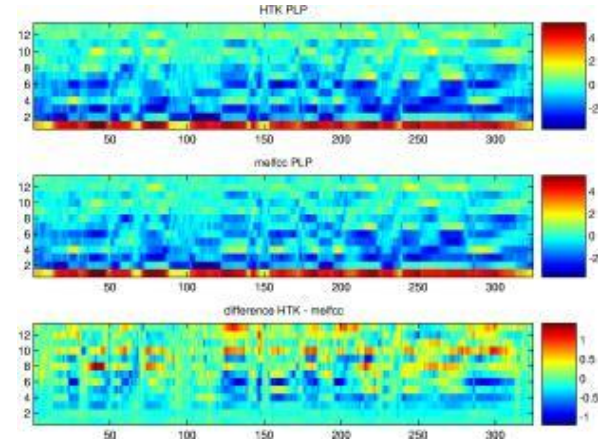
GIST



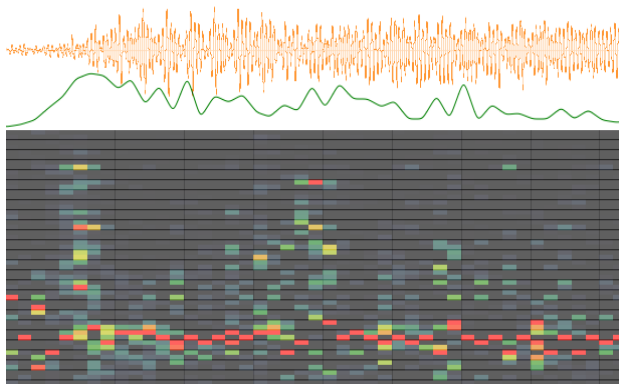
Old school: hand-craft features




Spectrogram




MFCC



Flux

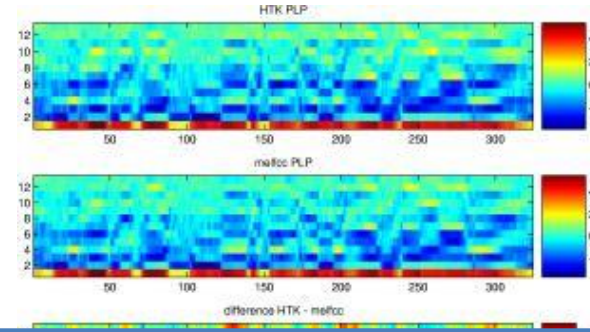
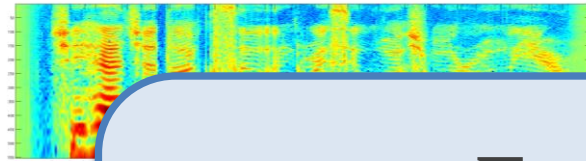
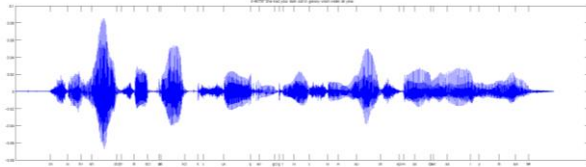


ZCR

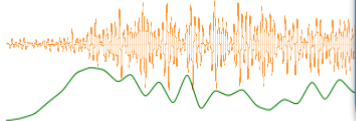


Rolloff


Old school: hand-craft features




Feature learning:
Can we automatically learn
useful features?



Flux

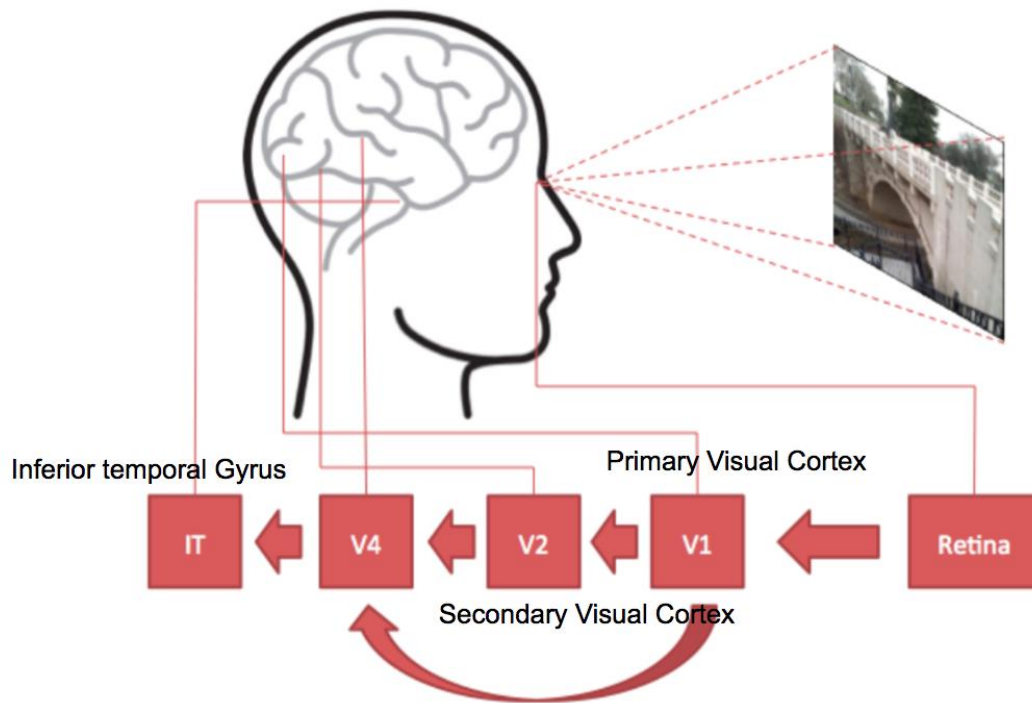


ZCR



Rolloff

Early inspirations from visual cortex



V1: Edge detection, etc.

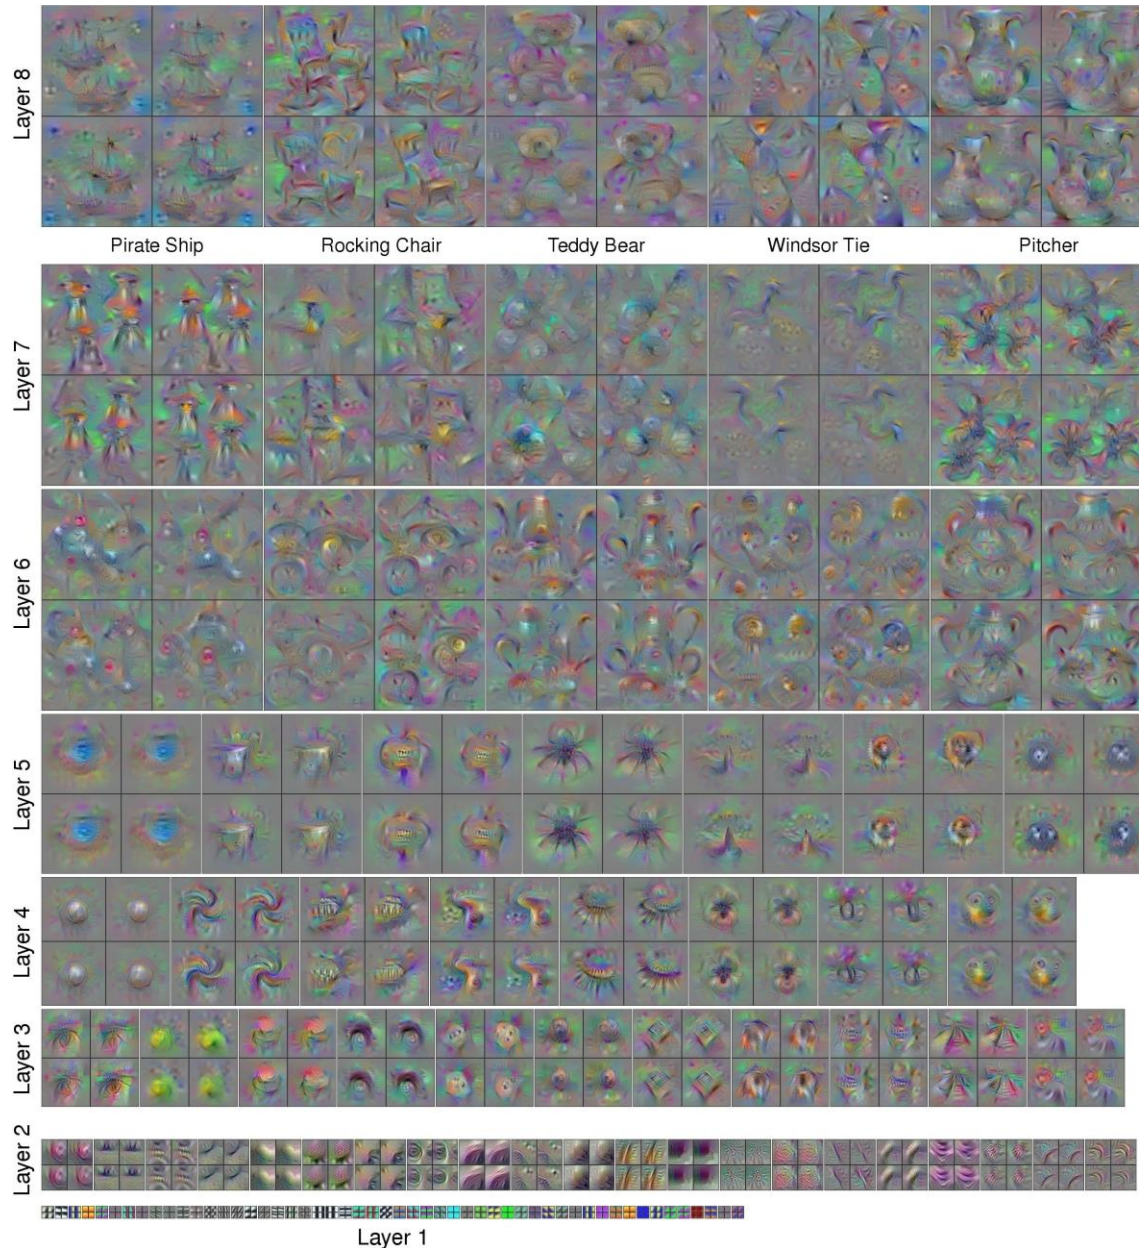
V2: Extract simple visual properties (orientation, spatial frequency, color, etc)

V4: Detect object features of intermediate complexity

TI: Object recognition.

Image: Wang, Raj, and Xing. [*"On the Origin of Deep Learning."*](#)

What do deep networks learn?



Yosinski et al '15:

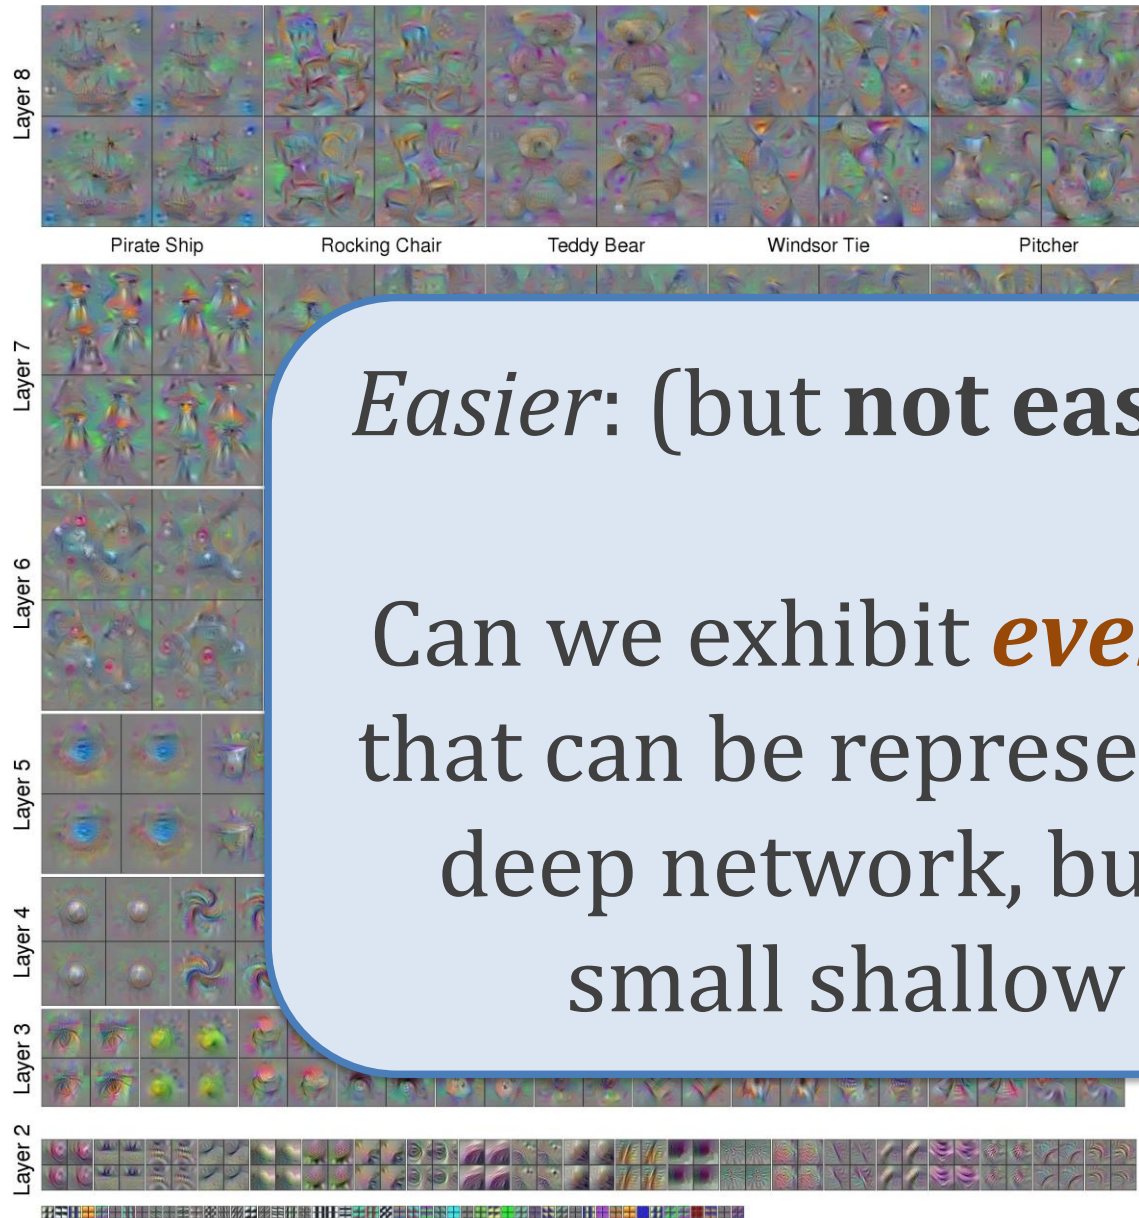
<http://yosinski.com/deepvis>

Q: What “patterns” do neurons respond to?

A: From random start, do gradient descent to find an input for which neuron activation* is high.

*: This produces completely unrecognizable images – they are regularized w/ an image prior.

What do deep networks learn?



Yosinski et al '15:

<http://yosinski.com/deepvis>

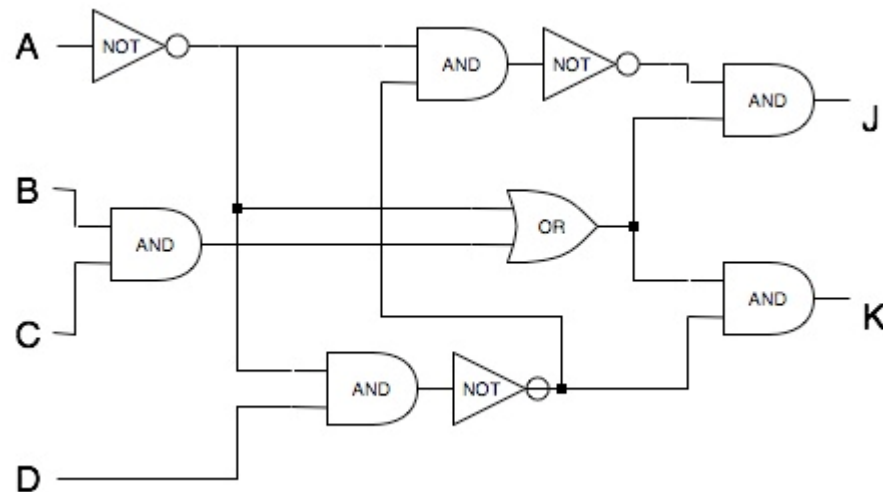
Easier: (but **not easy** it turns out!)

Can we exhibit *even one* function that can be represented as a small deep network, but cannot as a small shallow network?

A brief history of depth separation

Early roots: theoretical computer science

Boolean circuits: a directed acyclic graph model for computation over binary inputs; each node (“*gate*”) performs an operation (e.g. OR, AND, NOT) on the inputs from its predecessors.



A brief history of depth separation


Early roots: theoretical computer science

Boolean circuits: a directed acyclic graph model for computation over binary inputs; each node (“*gate*”) performs an operation (e.g. OR, AND, NOT) on the inputs from its predecessors.

Relaxation of the P vs NP question: separate more “structured” models of computation – e.g. shallow vs deep Boolean circuits.

Seminal result by *Håstad* ('86): **parity** function (calculates parity of number of ones in input) cannot be approximated by a small **constant-depth** circuit with OR and AND gates.

[Highly non-trivial; **Gödel Prize!!**]



Modern iterations of depth separation

Related architectures/models of computation

Sum-product networks [*Bengio, Delalleau '11*]

Weaker measures of complexity

Bound on number of linear regions for ReLU networks
[*Montufar, Pascanu, Cho, Bengio '14*]

True approximation error results

A small deep network cannot be approximated
by a smaller shallow network [*Telgarsky '15*]



Separating deep and shallow networks

Can be imitated for higher dim too

Family of functions for ever L

Theorem (Telgarsky '15): For every $L \in \mathbb{N}$,
there is a function $\mathbf{f}: [0,1] \rightarrow [0,1]$ representable as a network of
depth $O(L^2)$, with $O(L^2)$ nodes, and ReLU activation, s.t.
for every network $\mathbf{g}: [0,1] \rightarrow \mathbb{R}$ of depth L and $\leq 2^L$ nodes, and
ReLU activation, we have

$$\int_{[0,1]} |f(x) - g(x)| dx \geq \frac{1}{32}$$


*You cannot approximate f well
using any g , **even** in an
averaged (l_1) sense.*

*Outputting trivial
approximation $g = 0$ will give
error 1, so this is a really bad
approximation.*

Q: what can deep networks do easily?

A ReLU network f is **piecewise linear**: we can subdivide domain of f into a finite number of *polyhedral* pieces (P_1, P_2, \dots, P_N) , s.t. in each piece, f is linear. In other words, $\forall x \in P_i, f(x) = A_i x + b_i$.

(Once we know which ReLUs are in the linear regime, and which are zeroed out, the function f calculates is linear.)




Deeper networks can make a larger number of pieces.

Shallow functions have few linear pieces

Let's reason how the number of linear pieces behaves under compositions.

Claim: if $f: \mathbb{R} \rightarrow \mathbb{R}$ is a ReLU network with hidden layer widths (m_1, m_2, \dots, m_L) . Then, f has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ linear pieces.




Shallow functions have few linear pieces

Let's reason how the number of linear pieces behaves under compositions.

Claim: if $f: \mathbb{R} \rightarrow \mathbb{R}$ is a ReLU network with hidden layer widths (m_1, m_2, \dots, m_L) . Then, f has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ linear pieces.

Proof: Induction on L :

$L = 1$:



If only one hidden node: one value if $x \geq -\frac{b}{w}$
another if $x < -\frac{b}{w}$


Shallow functions have few linear pieces

Let's reason how the number of linear pieces behaves under compositions.

Claim: if $f: \mathbb{R} \rightarrow \mathbb{R}$ is a ReLU network with hidden layer widths (m_1, m_2, \dots, m_L) . Then, f has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ linear pieces.

Proof: Induction on L :

$L = 1$:



Each node introduces at most one breakpoint, e.g.:

$$\text{Let's assume } -\frac{b_2}{w_2} \geq -\frac{b_1}{w_2}$$

Different values in

$$\left[0, -\frac{b_1}{w_1}\right], \left[-\frac{b_1}{w_1}, -\frac{b_2}{w_2}\right], \left[-\frac{b_2}{w_2}, 1\right]$$


Shallow functions have few linear pieces

Let's reason how the number of linear pieces behaves under compositions.

Claim: if $f: \mathbb{R} \rightarrow \mathbb{R}$ is a ReLU network with hidden layer widths (m_1, m_2, \dots, m_L) . Then, f has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ linear pieces.

Proof: Induction on L :

$L = 1$:




Continuing, number of pieces is at most $m_1 + 1$, as we need.

Shallow functions have few linear pieces


Claim: if $f: \mathbb{R} \rightarrow \mathbb{R}$ is a ReLU network with hidden layer widths (m_1, m_2, \dots, m_L) . Then, f has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ linear pieces.

Proof: Inductive step:




Take any **node** in penultimate layer:

By inductive hypothesis, pre-activation of node is piecewise linear, $\leq 2^{L-2}(m_1 + 1) m_2 \dots m_{L-1}$ pieces.




Applying ReLU to pre-activation of **node** can at most double the number of pieces for that node !

Pre-activation
fn for **node**




Post-activation
fn for **node**



Shallow functions have few linear pieces


Claim: if $f: \mathbb{R} \rightarrow \mathbb{R}$ is a ReLU network with hidden layer widths (m_1, m_2, \dots, m_L) . Then, f has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ linear pieces.

Proof: Induction on L :



So, **output** is linear combination of m_L piecewise linear function, each with at most $2^{L-1}(m_1 + 1) m_2 \dots m_{L-1}$ pieces.

A linear comb. of 2 piecewise lin. fns, each w/ at most **a** and **b** pieces gives a piecewise lin. fn w/ at most **a+b** pieces:




Hence, output has at most $2^{L-1}(m_1 + 1) m_2 \dots m_L$ pieces.

Deep function with many oscillations

The player : the “triangle” map $\Delta: [0,1] \rightarrow \mathbb{R}$

$$\Delta(x) = 2\sigma(x) - 2\sigma(2x - 1) = \begin{cases} 2x, & \text{if } x \leq 1/2 \\ 2 - 2x, & \text{if } x > 1/2 \end{cases}$$


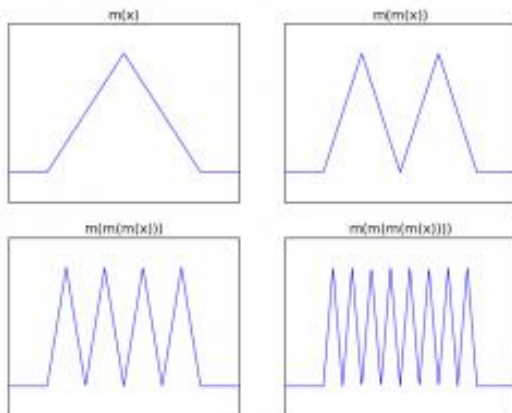


Deep function with many oscillations

$$\Delta(x) = 2\sigma(x) - 2\sigma(2x - 1) = \begin{cases} 2x, & \text{if } x \leq 1/2 \\ 2 - 2x, & \text{if } x > 1/2 \end{cases}$$

Idea: compose the triangle function w/ itself many times!

Idea: composing it k times should look like sawtooth with 2^{k-1} peaks




We'll show that function w/ small number of linear pieces can't approximate this well.

Deep function with many oscillations

$$\Delta(x) = 2\sigma(x) - 2\sigma(2x - 1) = \begin{cases} 2x, & \text{if } x \leq 1/2 \\ 2 - 2x, & \text{if } x > 1/2 \end{cases}$$

Claim: $\Delta^k(x) = \Delta\left(\underbrace{2^{k-1}x - \lfloor 2^{k-1}x \rfloor}_{\text{"Squish" triangle in every } 1/2^{k-1} \text{-sized interval}}\right)$

*"Squish" triangle in every
1/2^{k-1}-sized interval*



Proof: Induction:

K=1: by definition

K \Rightarrow K+1:


$$x \leq \frac{1}{2}: \quad \Delta^{k+1}(x) = \Delta^k(\Delta(x)) = \Delta^k(2x) = \Delta(2^k x - \lfloor 2^k x \rfloor)$$

Deep function with many oscillations

$$\Delta(x) = 2\sigma(x) - 2\sigma(2x - 1) = \begin{cases} 2x, & \text{if } x \leq 1/2 \\ 2 - 2x, & \text{if } x > 1/2 \end{cases}$$

Claim: $\Delta^k(x) = \Delta \left(\underbrace{2^{k-1}x - \lfloor 2^{k-1}x \rfloor}_{\text{“Squish” triangle in every } 1/2^{k-1} \text{-sized interval}} \right)$

*“Squish” triangle in every
1/2^{k-1}-sized interval*



Proof: Induction:

K=1: by definition

K \Rightarrow K+1:


$$\begin{aligned} x > \frac{1}{2}: \quad \Delta^{k+1}(x) &= \Delta^k(\Delta(x)) = \Delta^k(2 - 2x) = \Delta^{k-1}(\Delta(2 - 2x)) = \Delta^{k-1}(\Delta(1 - (2 - 2x))) \\ &= \Delta^{k-1}(\Delta(2x - 1)) = \Delta(2^k x - 2^k - \lfloor 2^k x - 2^k \rfloor) = \Delta(2^k x - \lfloor 2^k x \rfloor) \end{aligned}$$

Deep function with many oscillations

$$\Delta(x) = 2\sigma(x) - 2\sigma(2x - 1) = \begin{cases} 2x, & \text{if } x \leq 1/2 \\ 2 - 2x, & \text{if } x > 1/2 \end{cases}$$

Claim: $\Delta^k(x) = \Delta \left(\underbrace{2^{k-1}x - \lfloor 2^{k-1}x \rfloor}_{\text{“Squish” triangle in every } 1/2^{k-1} \text{ -sized interval}} \right)$


*“Squish” triangle in every
 $1/2^{k-1}$ -sized interval*



The function f in the theorem will be $\Delta^{2L^2+2}(x)$

We will show shallow g 's can't approximate f .


Shallow g's can't approximate f



Consider the line $y = \frac{1}{2}$. We will count (half) **triangles** of **f** on a different side than the **g**. Each contributes at least $\frac{1}{2} \times (\text{triangle area})$ to $\int |f - g| dx$:

$$\int_{x \in \text{triangle}} |f - g| dx \geq \int_{x \in \text{triangle}} \frac{1}{2} dx = \frac{1}{2} \times (\text{triangle area})$$

Shallow g's can't approximate f




Consider the line $y = \frac{1}{2}$. We will count (half) triangles of f on a different side than the g . Each contributes at least $\frac{1}{2} \times (\text{triangle area})$ to $\int |f - g| dx$.

Let $N_f = 2^{2L^2+2} - 1$ and N_g be the number of (half) triangles of f, g .

Since we get 2^{2L^2+1} copies of Δ , each gives 2 half triangles, but 1 is lost due to boundary

Shallow g's can't approximate f




Consider the line $y = 1/2$. We will count (half) triangles of f on a different side than the g . Each contributes at least $1/2 \times (\text{triangle area})$ to $\int |f - g| dx$.

Let $N_f = 2^{2L^2+2} - 1$ and N_g be the number of triangles of f, g .

For any **interval** of g corresponding to triangle: if there are m triangles of f contained **entirely** in interval, $\geq \frac{m-1}{2}$ lie on opposite side of g .

We “miss” at most N_g triangles (i.e. they are **not contained** entirely in an **interval** of g) – one for each boundary between triangles. Hence, we “contribute” at least $(N_f - N_g)/2$ triangles.

Shallow g's can't approximate f




Let N_f and N_g be the number of triangles of f , g . We “contribute” at least $(N_f - N_g)/2$ triangles.

Now, since shallow nets have small N_g : $N_g \leq (2n)^L \leq 2^{2L^2}$

We also showed that $N_f = 2^{2L^2+2} - 1$

\Rightarrow We “contribute” at least $2^{2L^2} - 1$ triangles.

Shallow g's can't approximate f



We “contribute” at least $2^{2L^2} - 1$ triangles.

Each triangle has area $\frac{1}{2} \times (\text{base}) \times (\text{height})$

$$= \frac{1}{2} \times \left(\frac{1}{2^{2L^2+2}} \right) \times \frac{1}{2} = \frac{1}{2^{2L^2+4}}$$

Total contribution: $\left(2^{2L^2} - 1 \right) \times \frac{1}{2^{2L^2+4}} \geq \frac{1}{32}$



Parting thoughts

Thm can be generalized to **d dimensions**, lots of other **activation functions**. (Need to bound # of pieces for multidim compositions.)

There can be benefits of depth to approximating **smooth** functions

Thm (Yarotsky '16). Suppose $f: [0, 1]^d \rightarrow \mathbb{R}$ has all partial derivs of order r coordinate-wise bdd in $[-1, +1]$ and let $\epsilon > 0$ be given. Then there exists a $O\left(\ln \frac{1}{\epsilon}\right)$ – depth and $\left(\frac{1}{\epsilon}\right)^{O\left(\frac{d}{r}\right)}$ – size network so that $\sup_{x \in [0, 1]^d} |f(x) - g(x)| \leq \epsilon$

Interaction of depth w/ **architecture**: depth tends to be problematic from an optimization point of view (“*vanishing gradient problem*”).

Lots of proposal architectures to remedy this (**ResNet, DenseNet**)

In general, interplay of depth with both **optimization** and **generalization** is **not** well understood.



Stay tuned!