

10707

Deep Learning: Spring 2020

Andrej Risteski


Machine Learning Department

Lecture 11:

Variational methods

Graphical Models

Recall: **graph** contains a set of nodes connected by edges.



In a **probabilistic graphical model**, each node represents a random variable, links represent “probabilistic dependencies” between random variables.



Graph specifies how joint distribution over all random variables **decomposes** into a **product** of factors, each factor depending on a subset of the variables.

Two types of graphical models:



- **Bayesian networks**, also known as **Directed Graphical Models** (the links have a particular directionality indicated by the arrows)
- **Markov Random Fields**, also known as **Undirected Graphical Models** (the links do not carry arrows and have no directional significance).

Algorithmic pros/cons of latent-variable models (so far)

RBM's

- ⌘ Hard to draw samples 
(In fact, #P-hard provably, even in Ising models)
- ⌘ Easy to sample posterior distribution over latents 

Directed models

- ⌘ Easy to draw samples 
- ⌘ Hard to sample posterior distribution over latents 
(In fact, #P-hard even in mixtures)

Canonical tasks with graphical models

Inference

Given values for the parameters θ of the model, *sample/calculate* marginals (e.g. sample $p_\theta(x_1)$, $p_\theta(x_4, x_5)$, $p_\theta(z|x)$, etc.)

Learning

Find values for the parameters θ of the model, that give a *high likelihood* for the observed data. (e.g. canonical way is solving maximum likelihood optimization

$$\max_{\theta \in \Theta} \sum_{i=1}^n \log p(x_i)$$

Other methods exist, e.g. method of moments (matching moments of model), but less used in deep learning practice.

Canonical tasks with graphical models

Inference

Inference is hard in undirected fully-observable models (due to **partition function**); easy in fully-observable Bayesian nets.

It's easy for RBM's, hard for latent-variable Bayesian nets (again, implicit **normalizing factor** is hard.)

Learning

We will derive “iterative”/“incremental” learning algorithms, using inference algorithms as subroutines.

We will, in particular see how a technique called “**variational methods**” can be used. (Next time, we see how **MCMC** methods can be used.)



Inference

Algorithmic approaches

When faced with a difficult to calculate probabilistic quantity (partition function, difficult posterior), there are two families of approaches:


MARKOV CHAIN MONTE CARLO

- ❖ **Random walk** w/ equilibrium distribution the one we are trying to sample from.
- ❖ Well studied in TCS.



VARIATIONAL METHODS

- ❖ Based on solving an **optimization** problem.
- ❖ Very popular in practice.
- ❖ Comparatively poorly understood



Part I: Inference in undirected graphical models

Though not always true, calculating marginals is often reducible to calculating partition functions.

Simple example: Ising models

$$P_{\theta}(\mathbf{x}) = \frac{1}{Z(\theta)} \exp \left(\sum_{ij \in E} x_i x_j \theta_{ij} + \sum_{i \in V} x_i \theta_i \right)$$

Partition fn of an appropriately modified Ising model

$$P_{\theta}(x_k = 1) = \frac{\sum_{\mathbf{x}_{-k} \in \{-1,1\}^{n-1}} \exp(\theta_k + \sum_{kj \in E} x_k \theta_j + \sum_{ij \in E, i \neq k, j \neq k} x_i x_j \theta_{ij} + \sum_{i \neq k} \theta_i x_i)}{\sum_{\mathbf{x} \in \{-1,1\}^n} \exp(\sum_{ij} x_i x_j \theta_{ij} + \sum_i \theta_i x_i)}$$

Partition fn of original Ising model

Part I: Inference in undirected graphical models

Though not always true, calculating marginals is often reducible to calculating partition functions.

Simple example: Ising models

$$P_{\theta}(\mathbf{x}) = \frac{1}{Z(\theta)} \exp \left(\sum_{ij \in E} x_i x_j \theta_{ij} + \sum_{i \in V} x_i \theta_i \right)$$

Formal term for when sampling reduces to calculating partition functions: *self-reducible problems*.

How do we calculate/approximate the partition function?


Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{p: \text{distribution over } \mathcal{X}} H(p) + \mathbb{E}_{x \sim p}[E(x)]$$

Find the distribution that has both high entropy, and high expected energy value

$$H(p) := - \sum_{x \in \mathcal{X}} p(x) \log p(x)$$



Variational methods for partition functions


Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Proof:
$$0 \leq KL(q || p) = \mathbb{E}_q \log q - \mathbb{E}_q \log p$$
$$= -H(q) - \mathbb{E}_{x \sim q}[E(x)] + \log Z$$

$$H(q) + \mathbb{E}_{x \sim q}[E(x)] \leq \log Z$$

Hence,
$$\max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)] \leq \log Z$$



Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$


Proof: $0 \leq KL(q \parallel p) = \mathbb{E}_q \log q - \mathbb{E}_q \log p$

$$= -H(q) - \mathbb{E}_{x \sim q}[E(x)] + \log Z$$

$$H(q) + \mathbb{E}_{x \sim q}[E(x)] \leq \log Z$$

Equality is attained if $p = q$: $KL(q \parallel p) = 0$, so

$$H(q) + \mathbb{E}_{x \sim q}[E(x)] = \log Z$$



Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Hence, we've reduced calculating partition function to an optimization problem!

But, there is a **serious issue**: how do we solve an optimization over the set of distributions over \mathcal{X} ?

Even if \mathcal{X} is a really simple domain, e.g. $\mathcal{X} = \{\pm 1\}^n$, the **trivial way** to solve the problem would involve introducing a variable $q(x)$, $\forall x \in \{\pm 1\}^n$: there are 2^n of them.

In fact, you can't be clever – there are results showing this can be #P hard even for Ising models!

Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

What can we do to try to approximate this expression?

Inspiration from physics: solve a *simpler* optimization problem over a *restricted class* of distributions we can explicitly parametrize.

Typical example: mean-field approximation

Consider again $\mathcal{X} = \{\pm 1\}^n$. A *product distribution* depends on n parameters only: since $p(x) = \prod_i p_i(x_i)$, for each $i \in [n]$, we only need to specify $p_i(x_i) = 1$.

Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Typical example: mean-field approximation

Consider again $\mathcal{X} = \{\pm 1\}^n$. A *product distribution* depends on n parameters only: since $q(x) = \prod_i q_i(x_i)$, for each $i \in [n]$, we only need to specify $q_i(x_i) = 1$.

Then, we will solve the optimization problem: $\max_{q = \prod_i q_i} H(q) + \mathbb{E}_{x \sim q}[E(x)]$

It's clear that the number of parameters is small at least.

If we can **take gradient wrt variables $q_i(x_i)$** we can at least do *gradient descent*. Objective in general is *non-convex* though, so technically, this **can fail !!** (But often works ok.)

Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Typical example: mean-field approximation: $\max_{q = \prod_i q_i} H(q) + \mathbb{E}_{x \sim q}[E(x)]$

Can we take gradients?

(1) *Entropy factorizes:* $H(q_1 q_2) = \sum_{x_1, x_2 \in \{\pm 1\}} q_1(x_1) q_2(x_2) \log(q_1(x_1) q_2(x_2))$

$$= \sum_{x_1, x_2} q_1(x_1) q_2(x_2) (\log q_1(x_1) + \log q_2(x_2)) = \sum_{x_2} q_2(x_2) H(q_1) + \sum_{x_1} q_1(x_1) H(q_2)$$

$$= H(q_1) + H(q_2) = q_1(x_1) \log q_1(x_1) + (1 - q_1(x_1)) \log(1 - q_1(x_1)) + H(q_2)$$

Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Typical example: mean-field approximation: $\max_{q = \prod_i q_i} H(q) + \mathbb{E}_{x \sim q}[E(x)]$

Can we solve this?

(2) $\mathbb{E}_{x \sim q}[E(x)]$ is often not a problem:

e.g. for Ising models, $E(x) = \sum_{ij} J_{ij} x_i x_j$ so $\mathbb{E}_q[E(x)] = \sum_{ij} J_{ij} \mathbb{E}[x_i] \mathbb{E}[x_j]$

$\mathbb{E}[x_j] = q_j(x_j)$, so taking a gradient is simple.

Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Typical example: mean-field approximation: $\max_{q = \prod_i q_i} H(q) + \mathbb{E}_{x \sim q}[E(x)]$

Can we solve this?

(2) $\mathbb{E}_{x \sim q}[E(x)]$ is often not a problem:

Even if gradients are not explicit, if $E(x)$ is a sum of terms $\phi_S(x_S)$, we can estimate the expectations, and estimate the gradient from that.

(i.e. use zeroth order optimization method.)

Variational methods for partition functions

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Typical example: Gaussians with mean/covariance μ, Σ as parameters.

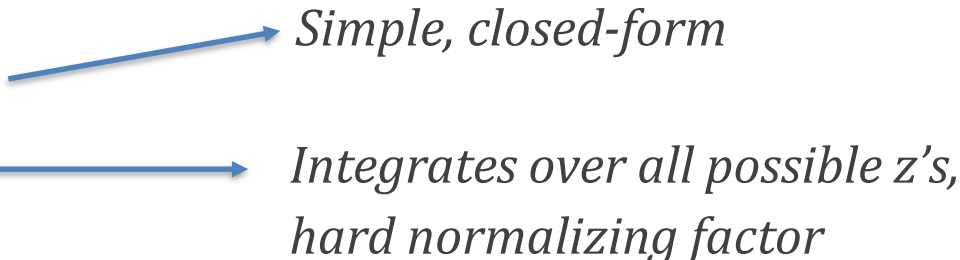
(1) Entropy of Gaussian has a closed-form formula:

$$\frac{1}{2} \log((2\pi e)^n \det \Sigma)$$

(2) Expectations wrt to Gaussian can be (often) estimated by drawing samples.

Part II: Inference in latent-variable Bayesian networks

As we noted last time, the difficulty for calculating posteriors in latent-variable models is of a *similar nature*:

$$p(z|x) = \frac{p(z)p(x|z)}{p(x)}$$


Simple, closed-form

*Integrates over all possible z 's,
hard normalizing factor*

Taking inspiration from previous part, we will approximate posteriors $p(z|x)$ in an analogous way:

Variational methods for posterior distributions

Let $p(z, x)$ be a joint distribution over latent variables and observables. Then, same as in the *Gibbs principle* calculation:

$$\begin{aligned} KL(q(z|x) || p(z|x)) &= \mathbb{E}_{q(z|x)} \log q(z|x) - \mathbb{E}_{q(z|x)} \log p(z|x) \\ &= -H(q(z|x)) - \mathbb{E}_{z \sim q}[p(z, x)] + \log p(x) \end{aligned}$$

Hence,

$$\operatorname{argmax}_{q(z|x)} KL(q(z|x) || p(z|x)) = \operatorname{argmax}\{ \mathbb{E}_{q(z|x)} \log q(z|x) - \mathbb{E}_{q(z|x)} \log p(z, x) \}$$

As in the undirected case, if q is a simple distribution (e.g. product distribution, Gaussian), we can optimize this using gradient descent.

Yet another mean field strategy: coordinate ascent

Consider updating a *single* coordinate of the mean-field distribution, that is keep $q_{-i}(z_i|x)$ fixed, and optimize for $q_i(z_i|x)$. Rewriting, we have:

$$KL(q(z|x) || p(z|x))$$

$$= \mathbb{E}_{q(z|x)} \log q(z|x) - \mathbb{E}_{q(z|x)} \log p(z, x)$$

$$= \sum_i \mathbb{E}_{q_i(z_i|x)} \log q_i(z_i|x) - \mathbb{E}_{q_i(z_i|x)} \left[\mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x) \right]$$

$$= \mathbb{E}_{q_i(z_i|x)} \log q_i(z_i|x) - \mathbb{E}_{q_i(z_i|x)} [\log \tilde{p}(z_i, x)] + C$$



Renormalize to make it a distribution

Yet another mean field strategy: coordinate ascent

Consider updating a *single* coordinate of the mean-field distribution, that is keep $q_{-i}(z_i|x)$ fixed, and optimize for $q_i(z_i|x)$. Rewriting, we have:


$$KL(q(z|x) || p(z|x))$$

$$= \mathbb{E}_{q(z|x)} \log q(z|x) - \mathbb{E}_{q(z|x)} \log p(z, x)$$

$$= \sum_i \mathbb{E}_{q_i(z_i|x)} \log q_i(z_i|x) - \mathbb{E}_{q_i(z_i|x)} \left[\mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x) \right]$$

$$= \mathbb{E}_{q_i(z_i|x)} \log q_i(z_i|x) - \mathbb{E}_{q_i(z_i|x)} [\log \tilde{p}(z_i, x)] + C$$

$$= KL(q_i(z_i|x) || \tilde{p}(z_i, x)) + C$$



Optimum is $q_i(z_i|x) = \tilde{p}(z_i, x)$

$$= \frac{\mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x)}{\int_{z_i} \mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x)}$$

Yet another mean field strategy: coordinate ascent

Consider updating a *single* coordinate of the mean-field distribution, that is keep $q_{-i}(z_i|x)$ fixed, and optimize for $q_i(z_i|x)$. Rewriting, we have:

$$KL(q(z|x) || p(z|x))$$

Iterate!

$$= \mathbb{E}_{q(z|x)} \log q(z|x) - \mathbb{E}_{q(z|x)} \log p(z, x)$$

$$= \sum_i \mathbb{E}_{q_i(z_i|x)} \log q_i(z_i|x) - \mathbb{E}_{q_i(z_i|x)} \left[\mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x) \right]$$


$$= \mathbb{E}_{q_i(z_i|x)} \log q_i(z_i|x) - \mathbb{E}_{q_i(z_i|x)} [\log \tilde{p}(z_i, x)] + C$$

$$= KL(q_i(z_i|x) || \tilde{p}(z_i, x)) + C$$


Optimum is $q_i(z_i|x) = \tilde{p}(z_i, x)$

$$= \frac{\mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x)}{\int_{z_i} \mathbb{E}_{q_{-i}(z_{-i}|x)} \log p(z_i, z_{-i}, x)}$$

What if we changed the order of p, q ?



Approximation is too compact.



Approximation is too spread.


What if we changed the order of p, q?

$$\text{KL}(q||p) = - \int q(\mathbf{Z}) \ln \frac{p(\mathbf{Z})}{q(\mathbf{Z})} d\mathbf{Z}.$$

There is a large positive contribution to the KL divergence from regions of \mathbf{Z} space in which:

- $p(\mathbf{Z})$ is near zero
- unless $q(\mathbf{Z})$ is also close to zero.

Minimizing $\text{KL}(q||p)$ leads to distributions $q(\mathbf{Z})$ that **avoid regions in which $p(\mathbf{Z})$ is small.**




What if we changed the order of p, q?

$$\text{KL}(p||q) = - \int p(\mathbf{Z}) \ln \frac{q(\mathbf{Z})}{p(\mathbf{Z})} d\mathbf{Z}.$$


There is a large positive contribution to the KL divergence from regions of \mathbf{Z} space in which:

- $q(\mathbf{Z})$ is near zero,
- unless $p(\mathbf{Z})$ is also close to zero.

Minimizing $\text{KL}(p||q)$ leads to distributions $q(\mathbf{Z})$ that **are nonzero in regions where $p(\mathbf{Z})$ is nonzero.**



Multimodal distributions



Blue contours show bimodal distribution, red contours single Gaussian distribution that best approximates it.

$KL(q||p)$ will tend to find a single mode, whereas $KL(p||q)$ will average across all of the modes.

Learning

Learning latent-variable directed graphical models

How should we try to learn the parameters of a graphical model?

The most obvious strategy: maximum likelihood estimation

Given data x_1, x_2, \dots, x_n , solve the optimization problem

$$\max_{\theta \in \Theta} \sum_{i=1}^n \log p(x_i)$$

Latent variables: we will use the variational principle again!

$$\log_{\theta} p(x) = \max_{q(z|x): \text{distribution over } \mathcal{Z}} H(q(z|x)) + \mathbb{E}_{q(z|x)}[\log p_{\theta}(x, z)]$$

Hence, MLE objective can be written as double maximization:

$$\max_{\theta \in \Theta} \max_{\{q_i(z|x_i)\}} \sum_{i=1}^n H(q_i(z|x_i)) + \mathbb{E}_{q_i(z|x_i)}[\log p_{\theta}(x_i, z)]$$

Variational methods for posterior distributions

ELBO (Evidence Lower Bound): Let $p(z, x)$ be a joint distribution over latent variables and observables. Then:

$$\log p(x) = \max_{q(z|x): \text{distribution over } \mathcal{Z}} H(q|z) + \mathbb{E}_{q(z|x)}[\log p(x, z)]$$

Write, by Bayes rule, $p(z|x) = \frac{p(x,z)}{p(x)}$. Then, the formula above follows by Gibbs variational principle with $E(x) = \log p(x, z)$. **Argmax = p(z|x)!**

Gibbs variational principle: Let $p(x) = \frac{1}{Z} \exp(E(x))$ be a distribution over a domain \mathcal{X} . Then, Z is the solution to the following optimization problem:

$$\log Z = \max_{q: \text{distribution over } \mathcal{X}} H(q) + \mathbb{E}_{x \sim q}[E(x)]$$

Expectation-maximization/ variational inference

The canonical algorithm for learning a single-layer latent-variable Bayesian network is an iterative algorithm as follows.

Consider the max-likelihood objective, rewritten as in the previous slide:

$$\max_{\theta \in \Theta} \max_{\{q_i(z|x_i)\}} \sum_{i=1}^n H(q_i(z|x_i)) + \mathbb{E}_{q_i(z|x_i)}[\log p_{\theta}(x_i, z)]$$

Algorithm maintains iterates $\theta^t, \{q_i^t(z|x_i)\}$, and updates them iteratively

(1) Expectation (E)-step:

Keep θ^t fixed, set $\{q_i^{t+1}(z|x_i)\}$, s.t. they maximize the objective above.

(2) Maximization (M)-step:

Keep $\{q_i^t(z|x_i)\}$ fixed, set θ^{t+1} s.t. it maximizes the objective above.

Clearly, every step cannot make the objective worse!

Does *not* mean it converges to global optimum – could, e.g. get stuck in a local minimum.

Expectation-maximization/ variational inference

The canonical algorithm for learning a single-layer latent-variable Bayesian network is an iterative algorithm as follows.

Consider the max-likelihood objective, rewritten as in the previous slide:

$$\max_{\theta \in \Theta} \max_{q_i(z|x_i)} \sum_{i=1}^n H(q_i(z|x_i)) + \mathbb{E}_{q_i(z|x_i)}[\log p_{\theta}(x_i, z)]$$

Algorithm maintains iterates $\theta^t, q_i^t(z|x_i)$, and updates them iteratively

(1) Expectation step:

Keep θ^t and set $q_i^{t+1}(z|x_i)$, s.t. they maximize the objective above.

If the class is infinitely rich, the optimum is $q_i^{t+1}(z|x_i) = p_{\theta^t}(z|x_i)$

This is called **expectation-maximization (EM)**.
If class is not infinitely rich, it's called **variational inference**.

Examples

1: Mixtures of spherical Gaussians

Consider a mixture of K Gaussians with unknown means $p = \sum_{i=1}^K \frac{1}{K} \mathcal{N}(\mu_i, I_d)$

Let's try to calculate the E and M steps.


E-step: the optimal $q_i^{t+1}(z|x_i)$ is $p_{\theta^t}(z|x_i)$. Can we calculate this?

By Bayes rule, $p_{\theta^t}(z = k|x_i) \propto p(x_i|z = k) \propto e^{-\|x_i - \mu_k^t\|^2}$

Writing out the normalizing constant, we have

$$p_{\theta^t}(z = k|x_i) = \frac{e^{-\|x_i - \mu_k^t\|^2}}{\sum_{k'} e^{-\|x_i - \mu_{k'}^t\|^2}}$$

*“Soft” version of assigning
point to nearest cluster*



Examples


1: Mixtures of spherical Gaussians

Consider a mixture of K Gaussians with unknown means $p = \sum_{i=1}^K \frac{1}{K} \mathcal{N}(\mu_i, I_d)$

Let's try to calculate the E and M steps.

M-step: given a guess $q_i^t(z|x_i)$, we can rewrite the maximization for θ as:

$$\max_{\theta \in \Theta} \sum_{i=1}^n H(q_i^t(z|x_i)) + \mathbb{E}_{q_i^t(z|x_i)} [\log p_{\theta}(x_i, z)]$$



$$= \mathbb{E}_{q_i^t(z|x_i)} [\log q_i^t(z|x_i) + \log p_{\theta}(x|z)]$$
$$\mathbb{E}_{q_i^t(z|x_i)} [\log p_{\theta}(x|z)]$$

Doesn't depend on θ

Examples

1: Mixtures of spherical Gaussians

Consider a mixture of K Gaussians with unknown means $p = \sum_{i=1}^K \frac{1}{K} \mathcal{N}(\mu_i, I_d)$

Let's try to calculate the E and M steps.

M-step: given a guess $q_i^t(z|x_i)$, we can rewrite the maximization for θ as:

$$\max_{\theta} \mathbb{E}_{q_i^t(z|x_i)} [\log p_{\theta}(x|z)] = \max_{\theta} - \sum_{i=1}^n \sum_{k=1}^K q_i^t(z = k|x_i) \|x_i - \mu_k\|^2$$

Setting the derivative wrt to μ_k to 0, we have:

$$\mu_k^{t+1} = \sum_{i=1}^n \frac{e^{-\|x_i - \mu_k^t\|^2}}{\sum_{k'} e^{-\|x_i - \mu_{k'}^t\|^2}} x_i$$

Examples

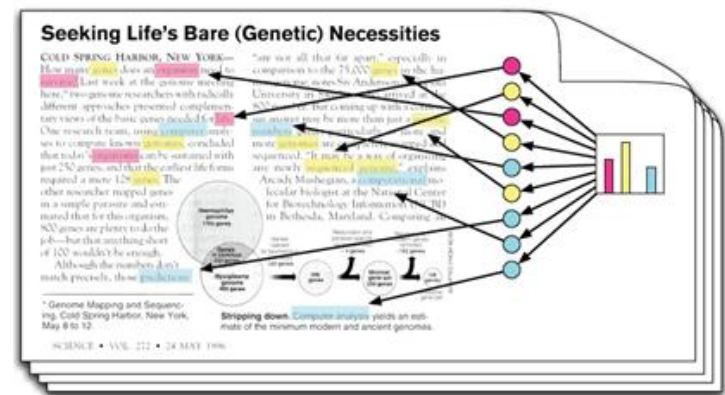

2: Latent Dirichlet Allocation

The **parameters** are: $\{\alpha_i\}_{i=1}^K$ (Dirichlet parameters) and **matrix** $\beta \in \mathbb{R}_+^{N \times K}$, where N is the size of the vocabulary.

The columns of β satisfy $\sum_{j=1}^N \beta_{ij} = 1$ (the **distribution of words** in a topic i)

To produce document:

- ❖ First, sample $\theta \sim \text{Dir}(\cdot | \alpha)$: this will be the **topic proportion vector** for the document.
- ❖ Each word in the document is generated in order, independently.
- ❖ To generate word i :
 - ❖ **Sample topic** z_i with categorical distribution with parameters θ
 - ❖ **Sample word** w_i with categorical distribution with parameters β_{z_i}



Examples

The E-step cannot be done in closed form:

$$p(\vec{\theta}_{1:D}, z_{1:D,1:N}, \vec{\beta}_{1:K} \mid w_{1:D,1:N}, \alpha, \eta) = \frac{p(\vec{\theta}_{1:D}, \vec{z}_{1:D}, \vec{\beta}_{1:K} \mid \vec{w}_{1:D}, \alpha, \eta)}{\int_{\vec{\beta}_{1:K}} \int_{\vec{\theta}_{1:D}} \sum_{\vec{z}} p(\vec{\theta}_{1:D}, \vec{z}_{1:D}, \vec{\beta}_{1:K} \mid \vec{w}_{1:D}, \alpha, \eta)}$$

(In fact, can be shown to be #P-hard to perform in the worst case.)

The variational family to approximate the posterior is commonly chosen to be a mean-field family:

$$q(\vec{\theta}_{1:D}, z_{1:D,1:N}, \vec{\beta}_{1:K}) = \prod_{k=1}^K q(\vec{\beta}_k \mid \vec{\lambda}_k) \prod_{d=1}^D \left(q(\vec{\theta}_{dd} \mid \vec{\gamma}_d) \prod_{n=1}^N q(z_{d,n} \mid \vec{\phi}_{d,n}) \right)$$

- **Probability of topic z given document d :** $q(\theta_d \mid \gamma_d)$
Each document has its own Dirichlet prior γ_d
- **Probability of word w given topic z :** $q(\beta_z \mid \lambda_z)$
Each topic has its own Dirichlet prior λ_z
- **Probability of topic assignment to word $w_{d,n}$:** $q(z_{d,n} \mid \phi_{d,n})$
Each word position $word[d][n]$ has its own prior $\phi_{d,n}$

Examples

$$q(\vec{\theta}_{1:D}, z_{1:D, 1:N}, \vec{\beta}_{1:K}) = \prod_{k=1}^K q(\vec{\beta}_k | \vec{\lambda}_k) \prod_{d=1}^D \left(q(\vec{\theta}_{dd} | \vec{\gamma}_d) \prod_{n=1}^N q(z_{d,n} | \vec{\phi}_{d,n}) \right)$$

- **Probability of topic z given document d :** $q(\theta_d | \gamma_d)$

Each document has its own Dirichlet prior γ_d

- **Probability of word w given topic z :** $q(\beta_z | \lambda_z)$

Each topic has its own Dirichlet prior λ_z

- **Probability of topic assignment to word $w_{d,n}$:** $q(z_{d,n} | \phi_{d,n})$

Each word position $word[d][n]$ has its own prior $\phi_{d,n}$

Parameter
updates:

One iteration of mean field variational inference for LDA

(1) For each topic k and term v :

$$(8) \quad \lambda_{k,v}^{(t+1)} = \eta + \sum_{d=1}^D \sum_{n=1}^N 1(w_{d,n} = v) \phi_{n,k}^{(t)}.$$

(2) For each document d :

(a) Update γ_d :

$$(9) \quad \gamma_{d,k}^{(t+1)} = \alpha_k + \sum_{n=1}^N \phi_{d,n,k}^{(t)}.$$

(b) For each word n , update $\phi_{d,n}$:

$$(10) \quad \phi_{d,n,k}^{(t+1)} \propto \exp \left\{ \Psi(\gamma_{d,k}^{(t+1)}) + \Psi(\lambda_{k,w_n}^{(t+1)}) - \Psi(\sum_{v=1}^V \lambda_{k,v}^{(t+1)}) \right\},$$

where Ψ is the digamma function, the first derivative of the $\log \Gamma$ function.

$$\beta_{ij} \propto \sum_{d=1}^M \sum_{n=1}^{N_d} \phi_{dni} w_{dn}^j.$$